

A METHOD FOR THE SPECTROSCOPIC INFERENCE OF FUNDAMENTAL STELLAR PARAMETERS

MAY 5, 2014

IAN CZEKALA¹ ET AL.

1. INTRODUCTION

- Where our technique fits in the ecosystem of stellar methods.
- Applicable to any kind of spectrum (long-slit, echelle, infrared, flux-calibrated or not)
- Examples of fields where accurate and unbiased stellar parameters are crucial. Exoplanets. T Tauri stars.

2. METHOD

Typically, the stellar properties we most desire are effective temperature T_{eff} , radius or surface gravity $\log g$, and metallicity $[\text{Fe}/\text{H}]$. Several high-quality libraries of synthetic spectra now exist (PHOENIX, Kurucz, more being developed for GAIA) and are parameterized by these fundamental stellar parameters, which together, we call $\vec{\theta}_{\star, \text{grid}}$. Synthetic libraries provide high-resolution spectra which span a range of $\vec{\theta}_{\star, \text{grid}}$ covering the main sequence of spectral types. The library is typically specified on a grid of equal spacing in T_{eff} , $\log g$, and $[\text{Fe}/\text{H}]$.

In addition to these fundamental parameters, a star also has several observed properties that are a function of its kinematics, geometry, and location in our galaxy: projected rotational velocity $v \sin i$, line of sight velocity v_z , total extinction A_V , and solid angle Ω . We call these secondary parameters $\vec{\theta}_{\star, \text{post}}$, because we can model their effects on the synthetic spectra in a “post-processing” step by convolution with an appropriate kernel or multiplication with a smooth function. Together, we call $\vec{\theta}_{\star} = \{\vec{\theta}_{\star, \text{grid}}, \vec{\theta}_{\star, \text{post}}\}$.

Using a synthetic stellar library as a backend and applying post-processing techniques, we attempt to forward-model the observed stellar spectra in order to determine the best-fit $\vec{\theta}_{\star}$. Because properly sampled spectra have correlated pixels, we present a framework which self-consistently models these correlations.

2.1. Interpolation from synthetic spectral libraries

Creating a new spectrum with a specific $\vec{\theta}_{\star, \text{grid}}$ requires either synthesizing a new spectrum using radiative transfer codes or interpolating from nearby grid points in the library. Because it is computationally expensive to synthesize a spectrum at high resolution over a wide wavelength range, we choose to interpolate. Spline interpolation for spectra (Husser 2012).

TODO: Implement band-limited interpolation from the synthetic grid using spline interpolation. Caution: As is, 100K/0.5 $\log g$ spacing in $\vec{\theta}_{\star, \text{grid}}$ may not be Nyquist sampling.

TODO: Investigate a “Bayesian Emulator” for using Gaussian processes for interpolating from a simulation grid.

2.2. Post-processing synthetic spectra to match the data

Libraries of synthetic spectra are generally computed at high resolution ($R \geq 100,000$), sampled with many pixels per resolution element, and do not include any rotational broadening or account for any instrumental effects. In order to transform a raw synthetic spectrum interpolated from the grid into one that matches the spectrum of a real star, we must post-process the spectrum to account for these secondary effects. The projected rotational velocity of the star, parameterized by $v \sin i$, broadens spectral line profiles and is mathematically described by convolution with a kernel $g_{v \sin i}(\lambda)$. UV, optical, and infrared spectra are acquired using a spectrograph which disperses light onto a CCD, defining a specific resolution and sampling rate for the spectrum. Because the resolution of the spectrograph and spacing of the pixels are different from the synthetic spectral library, we must convolve the raw spectrum with the line spread function (LSF) of the instrument $g_{\text{LSF}}(\lambda)$ and resample to the exact pixels of the CCD. In wavelength space, the $v \sin i$ and LSF operations are represented by the convolution of the synthetic spectrum with these two kernels

$$f_{\lambda, \text{inst}}(\lambda) = g_{v \sin i}(\lambda) \otimes g_{\text{LSF}}(\lambda) \otimes f_{\lambda, \text{synth}}(\lambda) \quad (1)$$

Using the convolution theorem, we can rewrite these operations as multiplications in Fourier space

$$F_{\lambda, \text{inst}}(s) = G_{\text{LSF}}(s) G_{v \sin i}(s) F_{\lambda, \text{synth}}(s) \quad (2)$$

where $G \leftrightarrow g$ and $F \leftrightarrow f$ are Fourier transform pairs. In order to use the fast Fourier transform (FFT) to execute these operations using discrete samples of the synthetic spectrum $f_{\lambda, \text{synth}}$, we must first resample $f_{\lambda, \text{synth}}$ so that it is on a uniform grid. In the case of a spectrum, the natural uniform grid is one that is equally sampled in velocity space, such there is an equal velocity shift Δv between pixels. This results in a wavelength grid that is linearly sampled in the logarithm of wavelength. Therefore, the Fourier coordinate s is the number of cycles per sampling interval, having units of inverse velocity [s/km]. Next, we multiply the Fourier transform of the synthetic spectrum by the Fourier transforms of the rotational velocity and line spread function kernels. Lastly, we do an inverse FFT to transform the modified spectrum $F_{\lambda, \text{inst}}$ back to wavelength space, where it is sampled at the wavelength locations of the pixels in the detector ($f_{\lambda, \text{inst}}$). When resampling the spectrum, we are careful to ensure band-limited interpolation using splines in order to prevent introducing non-physical high frequency structure into the spectrum.

Finally we apply corrections for interstellar extinction and Doppler shift. Interstellar extinction attenuates the spectrum by an amount A_{λ} , which is wavelength-dependent. We use the (name here) dust-extinction model, which is parameterized by

iczekala@cfa.harvard.edu

¹ Harvard-Smithsonian Center for Astrophysics, 60 Garden Street MS 10, Cambridge, MA 02138




Figure 1. **Top:** Chebyshev polynomial which modifies the model spectrum to account for inaccuracies in the flux-calibration. **Middle:** The data spectrum and model spectrum, after it has been interpolated, post-processed, and multiplied by the Chebyshev polynomial. **Bottom:** Residuals from the model fit. Note the large residual at $\lambda 5189\text{\AA}$ due to a missing opacity source from Ca I.

TODO: Model spectrum will be updated once I properly burn in the chain. Do these line styles look OK? I was going for something legible that might also print well in B/W.

extinction in the V photometric band, A_V . We correct for a radial velocity difference v_z between the star and the earth by Doppler shifting the model spectrum.

TODO: Doppler shifting the wavelengths could be implemented as a phase-shift in Fourier space, but it's pretty fast as-is.

TODO: We might try implementing the convolution in wavelength space, rather than the FFT. This would allow the LSF to change with wavelength.

Pixel width effects are not important as long as the LSF of the instrument is adequately sampled ($\gtrsim 3$ pixels across the LSF). If our spectrum were not properly sampled, we would need to additionally convolve with a boxcar the width of the pixels.

Due to small imperfections in the flux-calibration of the data spectrum, there may be broad regions of the spectrum which are offset from the true continuum level by a slight amount. A traditional approach to address this problem is to normalize both the data and the spectrum to the continuum level before comparison. While this may work well for hotter stars with well-defined continua, for cooler stars with large molecular features, normalization may introduce artificial features into the spectrum when the pseudo-continuum is incor-

rectly placed. Instead, we choose to multiply the model spectrum by a low-order Chebyshev polynomial that is normalized to unity. The coefficients of the polynomial, c_0, c_1, \dots , will be free-parameters in our model that we will solve for. If the behavior of the spectrograph is characterized by observations of spectrophotometric standards, then we can put priors on the coefficients that will prevent overfitting of the polynomials. For an example of a typical spectrum, and Chebyshev polynomial, see Figure 1.

With more uncertainty, the polynomial multiplication can also be used as a substitute for flux-calibration since the sensitivity function of an instrument is usually a smooth function with wavelength. When spectra are not flux calibrated, there is greater uncertainty about the properties of the spectrograph and wider priors must be used on the Chebyshev coefficients. In this case, the polynomial might destroy broad-scale information present in the spectrum by fitting it out, something that would be avoided with more accurate priors provided by the flux calibration for the model.

2.3. Likelihood function

To evaluate which parameters of the model M fit the data set D best we use a standard multidimensional Gaussian likelihood function which allows for covariance between data




Figure 2. The variety of residual behaviour, depending on the quality of model fit. From left to right: **Class 0** covariance results from slight model mismatch, and correlates nearby residuals. **Class I:** A missing absorption line in the model leaves a large, highly correlated patch of negative residuals. **Class II:** An extraneous line in the model leaves a large, highly correlated patch of positive residual. **Class III:** If lines are present in the model but of the wrong strength, many correlated residuals of moderate amplitude will result. The difficulty with class III lines is that for any specific line, there might exist a $\vec{\theta}_*$ that will fit the line, but there does not exist a $\vec{\theta}_*$ that will properly fit *all* the lines.

points. If the vector of residuals is

$$R(\vec{\theta}_*) = D - M(\vec{\theta}_*) \quad (3)$$

with a length of N data points, then the likelihood function is

$$p(D|\vec{\theta}_*) = \frac{1}{\sqrt{(2\pi)^N \det(C)}} \exp\left(-\frac{1}{2} R^T C^{-1} R\right) \quad (4)$$

and the log-likelihood function is

$$\ln[p(\vec{\theta}_*)] = -\frac{1}{2} R^T C^{-1} R - \frac{1}{2} \ln \det C - \frac{N}{2} \ln 2\pi \quad (5)$$

C is the covariance matrix which describes the covariance between pixels in the spectrum. For independent noise with noise per pixel σ_i , this matrix is diagonal

$$C = \begin{bmatrix} \sigma_0^2 & 0 & \cdots & 0 \\ 0 & \sigma_1^2 & \cdots & 0 \\ \vdots & & \ddots & \vdots \\ 0 & 0 & \cdots & \sigma_N^2 \end{bmatrix} \quad (6)$$

with $\sigma_{ij} = 0$ everywhere and equation 5 reduces to the familiar χ^2 form of a sum over the square of the residuals, weighted by the inverse variance of each data point

$$\ln[p(\vec{\theta}_*)] \propto -\frac{1}{2} \chi^2 = -\frac{1}{2} \sum_i \frac{R_i^2}{\sigma_i^2} \quad (7)$$

However, residuals from a spectroscopic fit are often correlated due to

TODO: DUE TO WHAT, exactly?? I'm confused. Wait for response from Tom Lored. Additionally, systematic errors

in the synthetic spectra (such as incorrect line strengths) will result in regions of highly correlated residuals. Properly accounting for correlations in the residuals requires that we use a covariance matrix with off-diagonal terms σ_{ij} and a likelihood function (Equation 5) which uses a matrix product instead of a sum over independent pixels.

We seek to understand and parameterize two types of covariance structure in our model, exemplified in Figure 2: the large scale but generally mild global covariance structure that results from [still to be finalized cause], and a regional covariance structure that is localized to areas where spectral line mismatch results in small patches of highly correlated residuals. Ignoring the complexity of spectroscopic residuals will bias the estimates of the model parameters $\vec{\theta}$.

As a simple analogy, consider the fit of a straight line to a dataset. If the noise in the dataset is correlated, then adjacent data points might be offset from the linear trend in the same direction by a similar amount. If the covariant structure of the noise is ignored, a simple χ^2 will treat these correlated offsets as part of the linear trend, which will result in a biased determination of the slope and intercept of the line, typically with uncertainties that are too small.

2.4. Parameterizing the covariance structure

We parameterize the covariance structure using covariance kernels, which sets the covariance between two pixels λ_i and λ_j . This behavior is analogous to the two-point correlation function used in cosmology, where the distance between two galaxies is used instead of wavelength.

Global covariance structure — A global covariance structure may result from a slight mismatch in the continuum or mild

mismatch of spectral lines like that shown as class 0 in Figure 2. To account for this structure, which is generally mild and typically exists across only a few pixels, we use a *stationary* covariance kernel, which satisfies the property that the degree of correlation only depends on the distance between the two pixels r . Functions like this are also called *radial basis functions*. For spectra, we map r to the velocity difference between two pixels

$$r(\lambda_i, \lambda_j) = \Delta v = \frac{c}{2} \left| \frac{\lambda_i - \lambda_j}{\lambda_i + \lambda_j} \right| \quad (8)$$

We choose the Matérn 3/2 kernel because...

TODO: Explore difference between using Matérn kernel and squared exponential kernel. Why Matérn and why Hann? Because they are the vanilla covariance functions that are tried and tested in the machine learning communities.

$$k_{\nu=3/2}(r, a, l) = a \left(1 + \frac{\sqrt{3}r}{l} \right) \exp \left(-\frac{\sqrt{3}r}{l} \right) \quad (9)$$

In order to keep the covariance matrix sparse for computational efficiency, we taper for compact support using a Hann window

$$w(r, r_0) = \frac{1}{2} + \frac{1}{2} \cos \left(\frac{\pi r}{r_0} \right) \quad (10)$$

for a global covariance kernel of

$$k_{\text{global}}(r, a, l, r_0) = w(r, r_0) k_{\nu=3/2}(r, a, l) \quad (11)$$

We add this covariance kernel to the independent variance already present along the matrix diagonal

$$k(\lambda_i, \lambda_j) = k_{\text{global}}(r, a, l, r_0) + \delta_{ij} \sigma_{ij}^2 \quad (12)$$

where δ_{ij} is the Kroenecker delta function.

To visualize what the covariance matrix C might look like when parameterized by Equation 11, see the left panel of Figure 3. To explore the relationship between the parameters of the covariance kernel and the properties of the noise, first consider simulating uncorrelated random Gaussian noise by drawing many independent samples from a Gaussian. This same uncorrelated noise could also be represented as a single draw from a N -dimensional Gaussian with a diagonal covariance matrix, like in Equation 6. If instead of a diagonal covariance matrix, we were to draw samples from a non-trivial covariance matrix like the one in the left panel of Figure 3, then we will see correlated noise like in the right panel.

TODO: Do you think examples like Figure 3 are helpful or a waste of space?

Specific line covariance — For the large and highly correlated residuals that result from class I, II, and III line errors, we use a *non-stationary* covariance kernel, which means that the covariance explicitly depends on λ_i and λ_j .

TODO: Explain how we come up with this kernel, which is the covariance resulting from a Gaussian function.

$$k_G(\lambda_i, \lambda_j | a, \lambda_\mu, \sigma) = \frac{a^2}{2\pi\sigma} \exp \left(-\frac{r^2(\lambda_i, \lambda_\mu) + r^2(\lambda_j, \lambda_\mu)}{2\sigma^2} \right) \quad (13)$$

The squared exponential kernel is

TODO: Why do we use the squared exponential and not

Matérn?

$k_{\text{exp}}(r, h) = \exp \left(\frac{-r^2}{2h^2} \right)$ (14)
 h is a bandwidth. If h is small, then there will be high-frequency structure. If h is large, then only low-frequency structure will remain. We also use the Hann window to taper the kernel for compact support

$$k(\lambda_i, \lambda_j | h, a, \mu, \sigma, r_0) = w(r, r_0) k_{\text{exp}}(r, h) k_G(\lambda_i, \lambda_j | a, \mu, \sigma) \quad (15)$$

TODO: Similar paragraph about visualizing the covariance matrix, drawing samples. References to Figure 4.

- non-stationary kernel localizes increased variance to a specific spectral line
- better than just masking out the region, since these contain info
- there will be many sets of regions (~1-4% of all lines might be covered by a region)
- there can be hyperparameters describing the *population* of regions in a hierarchical model (mostly width, amplitude).

Covariance kernel parameters are parameters in our model that we explicitly sample for.

The benefit mainly comes from simply modelling these systematic residuals to begin with, since a model with covariance is far more likely than forcing the fit, and far more justified than masking regions which do not fit.

TODO: Show how the kernel can accounts for complicated M dwarf structure where there is large, messy mismatch.

2.5. Exploring the posterior

Markov Chain Monte Carlo

TODO: Given this is 2014 and not 2010, how much do you think I need to explain?

2.6. Applications

Learnt covariance structure can be used to correct the models.

3. TESTS

TODO: Which tests to show?

This version of the paper was generated from a git repository available at <http://github.com/iancze/StellarSpectra/> with git hash bd9d54e (2014-05-05).

REFERENCES

Husser, T.-O. 2012, 3D-Spectroscopy of Dense Stellar Populations (Universitätsverlag Göttingen)




Figure 3. **Left** a covariance matrix generated with the Matérn kernel and typical parameters for our dataset. **Right** general spectroscopic residuals near a stellar continuum overlaid with a sample draw using the covariance matrix, showing that the two look similar in structure and amplitude.
 TODO: Prettify these figures. Consistent labelling. Change color stretch so Matrix zeros are white.




Figure 4. **Left** a covariance matrix generated with the region kernel. **Right** general spectroscopic residuals near a mismatched stellar line overlaid with a sample draw using the covariance matrix, showing that the two look similar in structure and amplitude.
 TODO: Similar improvements to Figure 3.